



BITS Pilani
Pilani Campus

COMPUTER ORGANIZATION AND SOFTWARE SYSTEMS

SESSION 11

Prof. C R Sarma
WILP.BITS-Pilani

Last Session



Contact Hour	List of Topic Title	Text/Ref Book/external resource
19-20	<ul style="list-style-type: none">Scheduling Algorithms FCFS, SJF, SRTF, Priority and RR	T2

Today's Session



Contact Hour	List of Topic Title	Text/Ref Book/external resource
21-22	<ul style="list-style-type: none">• Scheduling algorithms• Process Coordination	T2

Multilevel Queuing



- Process classification based on response-time requirements and scheduling needs
 - Foreground (interactive) processes
 - background (batch) processes
- Foreground processes may have priority (externally defined) over background processes
- Multilevel queue scheduling algorithm
 - Partition the ready queue into several separate queues

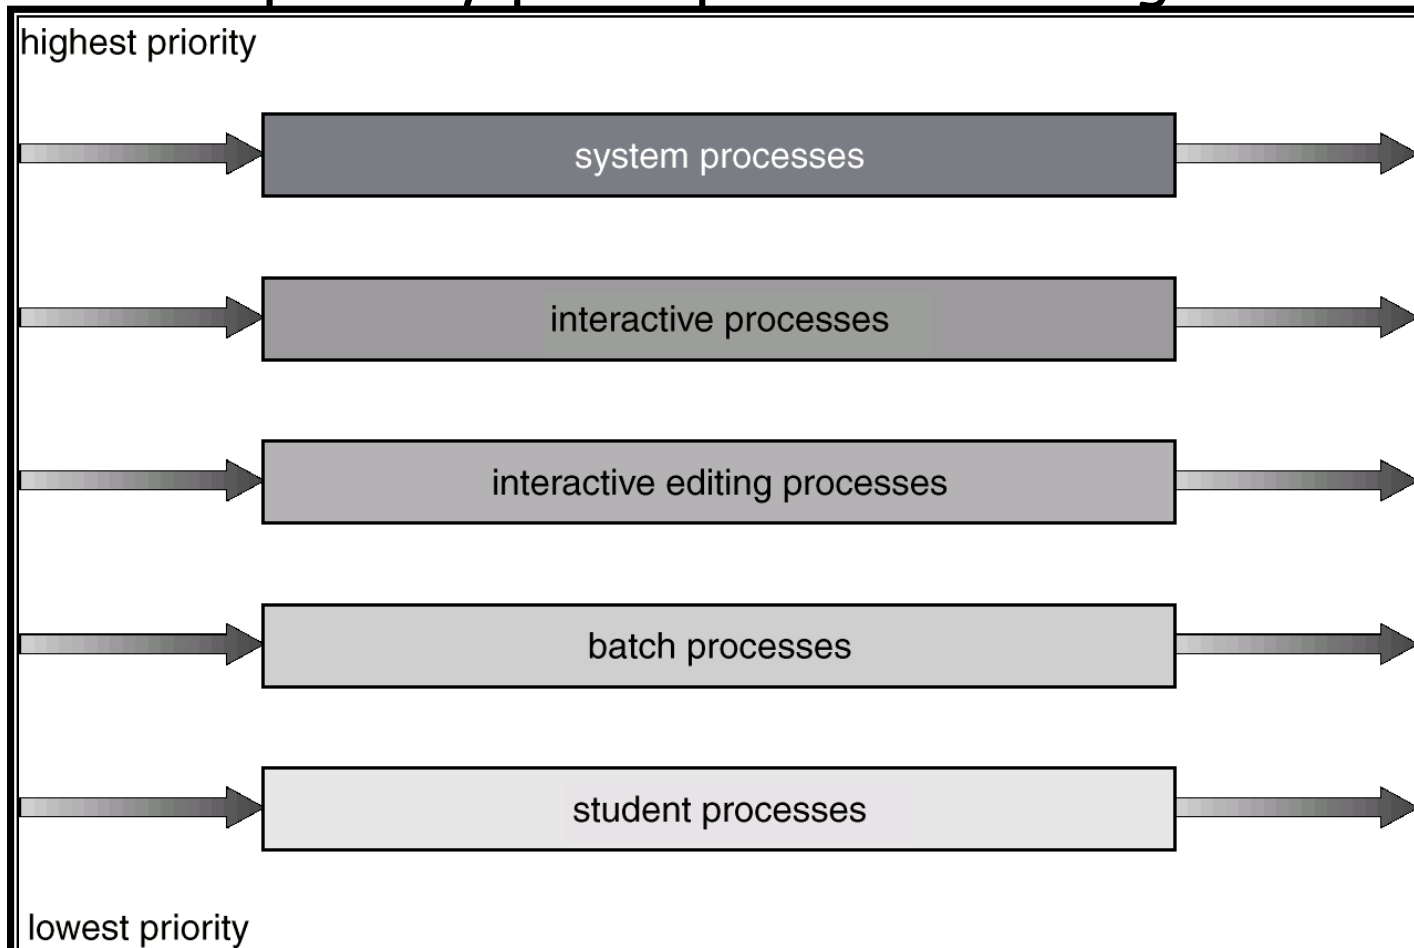
Contd...



- Process assignment to queue based on
 - Memory size, priority of process or process type
- Each queue has its own scheduling algorithm
- Example: foreground and background processes.
 - The foreground queue scheduled by an RR algorithm
 - background queue is scheduled by an FCFS algorithm.

Multilevel Queue Scheduling

- There must be scheduling among the queues
 - fixed priority preemptive scheduling



Multilevel Feedback Queue


- Disadvantage of Multilevel queue : Inflexible
- A process can move between the various queues; aging can be implemented this way.
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service

Example of Multilevel Feedback Queue



- Three queues:
 - Q_0 - time quantum 8 milliseconds
 - Q_1 - time quantum 16 milliseconds
 - Q_2 - FCFS
- Scheduling
 - A new job enters queue Q_0 . When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q_1 .
 - At Q_1 job receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q_2 .

Multilevel Feedback Queues



Multi-level Q

(Q1 & Q2 use RR, $P(Q1) > P(Q2)$, $tq1=4, tq2=3$)

Process	AT	BT	Priori ty	CT	TAT	WT
P1	0	10	2			
P2	3	7	1 (H)			
P3	4	6	2 (L)			
P4	12	5	1			
P5	18	8	1			

Q1									
Q2									



Process Synchronization

BITS Pilani
Pilani Campus

Introduction




- How do we maximize CPU utilization / improve efficiency?
 - Multiprogramming Vs Multiprocessing
- How to achieve concurrent operation in
 - Uniprocessor System Vs Multiprocessor system

uniprocessor system



multiprocessor system



Contd...



- Interleaving and overlapping improves processing efficiency, but may produce unpredictable results if not controlled properly
- Example:

```
Procedure echo;  
  Var out,in:Character;  
Begin  
  input (in, keyboard);  
  out:=in;  
  output(out,Display)  
End
```

Example (Contd...)



Process P1

1. input (in, keyboard);
2. -----
3. -----
4. -----
5. out:=in;
6. output(out, Display)

Process P2

1. -----
2. input (in, keyboard);
3. out:=in;
4. output(out, Display)

- Reasons unpredictable results :
 - Finite resources
 - Relative speed of execution of processes can not be predicted
 - Sharing of resources (non shareable) among processes
- Processes compete for resources
 - Deadlock
 - Mutual exclusion
 - Starvation
- Cooperating & competing processes can cause problem when executed concurrently → Synchronization

Process Synchronization



- Process **Synchronization** means sharing system resources by processes in a such a way that, concurrent access to shared data is handled thereby minimizing the chance of inconsistent data
- A **Critical Section** is a code segment that accesses shared variables or resources and has to be executed as an atomic action.
- Successful use of concurrency requires -
 - Ability to define critical section and
 - Enforce mutual exclusion

Process structure



Process P_i

do {

:

ENTRY SECTION

critical section

Exit SECTION

remainder section

} while (TRUE);

Main requirements



Three requirements

- **A mutual exclusion** : When one process is using a shared modifiable data, the other processes will be excluded from doing the same thing
- **Progress** : when no process in critical section, any process that makes a request is allowed to enter critical section without any delay
- **Bounded Waiting** : Processes requesting critical section should not be delayed indefinitely (no deadlock, starvation)

- ❖ *No assumption should be made about relative execution speed of processes or number of processes*
- ❖ *A process remains inside critical section for a finite amount of time*

Approach to handle Mutual Exclusion



- Software Approach (User is responsible for enforcing Mutual exclusion)
- Hardware Support - Disabling of Interrupt and using Special Instructions
- OS support - Semaphore

Critical Section (Solution1)

innovate

```
Procedure echo;  
  Var out,in:Character;  
Begin  
  input (in, keyboard);  
  out:=in;  
  output(out,Display)  
End
```

Process 0

```
while turn == 1 do {nothing }  
<Critical Section>  
turn = 1
```

Process 1

```
while turn == 0 do {nothing }  
<Critical Section>  
turn = 0
```

"I finished with it, now you have it" → Dekker's Algorithm

- Drawback 1: processes must strictly alternate - Pace of execution of one process is determined by pace of execution of other processes
- Drawback 2: if one processes fails other process is permanently blocked

Critical Section (Solution 2)

Peterson's solution



- Good algorithmic description of solving the problem
- Two process solution
- The two processes share two variables:
 - **int turn;**
 - **Boolean flag[2]**
- The variable **turn** indicates whose turn it is to enter the critical section
- The **flag** array is used to indicate if a process is ready to enter the critical section. **flag[i] = true** implies that process P_i is ready!

Algorithm



```
Process Pi
do {
    flag[i] = true;
    turn = j;
    while (flag[j] && turn == j);

    critical section

    flag[i] = false;
    -----
} while (true);
```

Algorithm



t	P0	P1
1	flag[0] = true; turn = 1;	
2		flag[1] = true; turn = 0;
3	while (flag[1] && turn == 1); <u>critical section</u>	
4		while (flag[0] && turn == 0); critical section
5	flag[0] = false	
6		while (flag[0] && turn == 0); critical section flag[1] = false;

Synchronization - Hardware Approach



- Protecting critical regions via locks

```
do {  
    acquire lock  
        critical section  
    release lock  
        remainder section  
} while (TRUE);
```


Synchronization - Hardware Approach...



- Uniprocessors - could disable interrupts
 - Currently running code would execute without preemption
 - Generally too inefficient on multiprocessor systems
 - Operating systems using this not broadly scalable
- Modern machines provide special atomic hardware instructions
 - **Atomic** = non-interruptible
 - Example : TestAndSet instruction, Swap instruction

Hardware approach - TestAndSet



- **TestAndSet** instruction used to write to a memory location and return its old value as a single atomic (i.e., non-interruptible) operation.
- **Definition of the TestAndSet () instruction:**

```
boolean TestAndSet(boolean *target) {  
    boolean rv = *target;  
    *target = true;  
    return rv;  
}
```

Mutual Exclusion with TestAndSet instruction



Shared data: Boolean lock = FALSE;

Process P_0

```
do {  
  while (TestAndSet( &lock))  
    ; //wait  
  critical section  
  lock = FALSE;  
  remainder section  
} while (TRUE);
```

Process P_1

```
do {  
  while (TestAndSet( &lock))  
    ; //wait  
  critical section  
  lock = FALSE;  
  remainder section  
} while (TRUE);
```

```
boolean TestAndSet(boolean *lock)  
{ boolean rv = *lock; *lock = TRUE; return rv; }
```

Semaphores

- is a variable which is treated in a special way
- allow processes to make use of critical section in exclusion to other processes
- process wanting to access critical section locks semaphore and releases lock on exit.
- is a synchronization tool

- Basic properties of semaphore:
 - semaphore S - integer variable with non negative values
 - can only be accessed via two operations
- wait (S):*
- ```
while $S \leq 0$ do no-op;
 $S--$;
```
- signal (S):*
- ```
 $S++$ ;
```
- Semaphore operation is atomic and indivisible
 - wait and signal operations are carried out without interruption

Types of semaphore



- Binary semaphore : can have two values 0 and 1
 - On some systems, binary semaphores are known as **mutex locks**, as they are locks that provide *mutual exclusion*.
- *Counting Semaphore* : integer value can range over an unrestricted domain.

Critical Section of n Processes



Shared data:

semaphore S ; *//initially $S = 1$*

Process P_i :

do {

 remainder section

wait(S);

 critical section

signal(S);

 remainder section

} while (1);

wait (S):

 while $S \leq 0$ do *no-op*;

$S--$;

signal (S):

$S++$;

Process synchronization



- Semaphore can be used to solve various synchronization problems
- Example : two concurrent processes : P1 (with S1) and P2 (with S2)
 - requirement : **s2 should be executed only after s1 is executed**
 - semaphore variable : **synch** initialized to zero.

P1:

S1;

signal (synch);

P2:

wait (synch);

S2;

wait (synch):

while $\text{synch} \leq 0$ do *no-op*;

synch--;

signal (synch):

synch++;

- Main disadvantage : Busy Waiting → waiting wastes CPU cycles
- semaphore is also called a **spinlock** because the process "spins" while waiting for the lock.
- Solution: on finding zero semaphore value (binary semaphore), the process can block itself instead of busy waiting
- The block operation places a process into a waiting queue associated with the semaphore, and the state of the process is switched to the waiting state.
- Then control is transferred to the CPU scheduler, which selects another process to execute.

- A process that is blocked, waiting on a semaphore S , should be restarted when some other process executes a `signal()` operation.
- The process is restarted by a `wakeup ()` operation, which changes the process from the waiting state to the ready state.
- The process is then placed in the ready queue.
- We need to modify the definition of the `wait ()` and `signal ()` semaphore operations.

Consumer – Producer Problem

- Also known as bounded buffer problem
 - Two processes – consumer and producer
 - Consumer and Producer processes share a common, fixed size buffer
 - Producer process : generates data and puts it in the buffer
 - Consumer process: consumes data from the buffer
 - *Problem statement : When a producer is placing an item in the buffer, then at the same time consumer should not consume any item.*
-

`mutex = 1`

`Full = 0` → Initially, all slots are empty.

`Empty = n` → All slots are empty



Producer

```
do{
//produce an item
wait(empty);
wait(mutex);
    //place in buffer
signal(mutex);
signal(full);

}while(true)
```

Consumer

```
do{
wait(full);
wait(mutex);
    // remove item from buffer
signal(mutex);
signal(empty);
    // consumes item

}while(true)
```

That's All for CS11

See Ya in

CS# 12

Multi-level Q

(Q1 & Q2 use RR, $P(Q1) > P(Q2)$, $tq1=4, tq2=3$)

Process	AT	BT	Priori ty	CT	TAT	WT	R _s
P1	0	10 7 4 1	2	36	36	26	0
P2 ✓	3	7 3 0	1 (H)	10	7	0	0
P3	4	6 2 0	2 (L)	35	31	25	21
P4 ✓	12	5 1 0	1	19	6	1	1
P5 ✓	18	8 4 0	1	26	8	0	0

$tq=4$	Q1	P2	P4	P5	P3			
$tq=3$	Q2	P1	P3	P1	P3			

P1	P2	P2	P1	P4	P4	P5	P5	P3	P1	P3	P1
0	3	7	10	13	17	18	22	26	29	32	36